

```

!python -m pip install
'git+https://github.com/facebookresearch/detectron2.git'

Collecting git+https://github.com/facebookresearch/detectron2.git
  Cloning https://github.com/facebookresearch/detectron2.git to
/tmp/pip-req-build-9ouujjrw
  Running command git clone --filter=blob:none --quiet
https://github.com/facebookresearch/detectron2.git /tmp/pip-req-build-
9ouujjrw
  Resolved https://github.com/facebookresearch/detectron2.git to
commit 57bdb21249d5418c130d54e2ebdc94dda7a4c01a
  Preparing metadata (setup.py) ... ent already satisfied: Pillow>=7.1
in /usr/local/lib/python3.10/dist-packages (from detectron2==0.6)
(8.4.0)
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.10/dist-packages (from detectron2==0.6) (3.7.1)
Requirement already satisfied: pycocotools>=2.0.2 in
/usr/local/lib/python3.10/dist-packages (from detectron2==0.6) (2.0.6)
Requirement already satisfied: termcolor>=1.1 in
/usr/local/lib/python3.10/dist-packages (from detectron2==0.6) (2.3.0)
Collecting yacs>=0.1.8 (from detectron2==0.6)
  Downloading yacs-0.1.8-py3-none-any.whl (14 kB)
Requirement already satisfied: tabulate in
/usr/local/lib/python3.10/dist-packages (from detectron2==0.6) (0.9.0)
Requirement already satisfied: cloudpickle in
/usr/local/lib/python3.10/dist-packages (from detectron2==0.6) (2.2.1)
Requirement already satisfied: tqdm>4.29.0 in
/usr/local/lib/python3.10/dist-packages (from detectron2==0.6)
(4.65.0)
Requirement already satisfied: tensorboard in
/usr/local/lib/python3.10/dist-packages (from detectron2==0.6)
(2.12.3)
Collecting fvcore<0.1.6,>=0.1.5 (from detectron2==0.6)
  Downloading fvcore-0.1.5.post20221221.tar.gz (50 kB)
  _____ 50.2/50.2 kB 1.2 MB/s eta
0:00:00
etadate (setup.py) ... detectron2==0.6)
  Downloading iopath-0.1.9-py3-none-any.whl (27 kB)
Collecting omegaconf>=2.1 (from detectron2==0.6)
  Downloading omegaconf-2.3.0-py3-none-any.whl (79 kB)
  _____ 79.5/79.5 kB 9.5 MB/s eta
0:00:00
detectron2==0.6)
  Downloading hydra_core-1.3.2-py3-none-any.whl (154 kB)
  _____ 154.5/154.5 kB 16.8 MB/s eta
0:00:00
detectron2==0.6)
  Downloading black-23.7.0-cp310-cp310-
manylinux2014_x86_64.whl (1.7 MB)
  _____ 1.7/1.7 MB 29.2 MB/s eta

```

0:00:00
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from detectron2==0.6) (23.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from fvcore<0.1.6,>=0.1.5->detectron2==0.6) (1.22.4)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from fvcore<0.1.6,>=0.1.5->detectron2==0.6) (6.0.1)
Collecting antlr4-python3-runtime==4.9.* (from hydra-core>=1.1->detectron2==0.6)
 Downloading antlr4-python3-runtime-4.9.3.tar.gz (117 kB)
117.0/117.0 kB 11.5 MB/s eta

0:00:00
etaddata (setup.py) ... iopath<0.1.10,>=0.1.7->detectron2==0.6)
 Downloading portalocker-2.7.0-py2.py3-none-any.whl (15 kB)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->detectron2==0.6) (1.1.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->detectron2==0.6) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->detectron2==0.6) (4.41.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->detectron2==0.6) (1.4.4)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->detectron2==0.6) (3.1.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->detectron2==0.6) (2.8.2)
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.10/dist-packages (from black->detectron2==0.6) (8.1.6)
Collecting mypy_extensions>=0.4.3 (from black->detectron2==0.6)
 Downloading mypy_extensions-1.0.0-py3-none-any.whl (4.7 kB)
Collecting pathspec>=0.9.0 (from black->detectron2==0.6)
 Downloading pathspec-0.11.1-py3-none-any.whl (29 kB)
Requirement already satisfied: platformdirs>=2 in /usr/local/lib/python3.10/dist-packages (from black->detectron2==0.6) (3.9.1)
Requirement already satisfied: tomli>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from black->detectron2==0.6) (2.0.1)
Requirement already satisfied: absl-py>=0.4 in /usr/local/lib/python3.10/dist-packages (from tensorboard-

```
>detectron2==0.6) (1.4.0)
Requirement already satisfied: grpcio>=1.48.2 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (1.56.0)
Requirement already satisfied: google-auth<3,>=1.6.3 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (2.17.3)
Requirement already satisfied: google-auth-oauthlib<1.1,>=0.5 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (1.0.0)
Requirement already satisfied: markdown>=2.6.8 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (3.4.3)
Requirement already satisfied: protobuf>=3.19.6 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (3.20.3)
Requirement already satisfied: requests<3,>=2.21.0 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (2.27.1)
Requirement already satisfied: setuptools>=41.0.0 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (67.7.2)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0
in /usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (0.7.1)
Requirement already satisfied: werkzeug>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (2.3.6)
Requirement already satisfied: wheel>=0.26 in
/usr/local/lib/python3.10/dist-packages (from tensorboard-
>detectron2==0.6) (0.40.0)
Requirement already satisfied: cachetools<6.0,>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3-
>tensorboard->detectron2==0.6) (5.3.1)
Requirement already satisfied: pyasn1-modules>=0.2.1 in
/usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3-
>tensorboard->detectron2==0.6) (0.3.0)
Requirement already satisfied: six>=1.9.0 in
/usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3-
>tensorboard->detectron2==0.6) (1.16.0)
Requirement already satisfied: rsa<5,>=3.1.4 in
/usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3-
>tensorboard->detectron2==0.6) (4.9)
Requirement already satisfied: requests-oauthlib>=0.7.0 in
/usr/local/lib/python3.10/dist-packages (from google-auth-
oauthlib<1.1,>=0.5->tensorboard->detectron2==0.6) (1.3.1)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
>tensorboard->detectron2==0.6) (1.26.16)
```

```
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
>tensorboard->detectron2==0.6) (2023.5.7)
Requirement already satisfied: charset-normalizer~=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
>tensorboard->detectron2==0.6) (2.0.12)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0-
>tensorboard->detectron2==0.6) (3.4)
Requirement already satisfied: MarkupSafe>=2.1.1 in
/usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1-
>tensorboard->detectron2==0.6) (2.1.3)
Requirement already satisfied: pyasn1<0.6.0,>=0.4.6 in
/usr/local/lib/python3.10/dist-packages (from pyasn1-modules>=0.2.1-
>google-auth<3,>=1.6.3->tensorboard->detectron2==0.6) (0.5.0)
Requirement already satisfied: oauthlib>=3.0.0 in
/usr/local/lib/python3.10/dist-packages (from requests-
oauthlib>=0.7.0->google-auth-oauthlib<1.1,>=0.5->tensorboard-
>detectron2==0.6) (3.2.2)
Building wheels for collected packages: detectron2, fvcore, antlr4-
python3-runtime
  Building wheel for detectron2 (setup.py) ... e=detectron2-0.6-cp310-
cp310-linux_x86_64.whl size=6114338
sha256=a976f533f006e201800add615bf643d37ff5d549f91bc646ae20bb1ec78ed43
8
  Stored in directory:
/tmp/pip-ephem-wheel-cache-95vy96fp/wheels/47/e5/15/94c80df2ba85500c5d
76599cc307c0a7079d0e221bb6fc4375
  Building wheel for fvcore (setup.py) ... e=fvcore-
0.1.5.post20221221-py3-none-any.whl size=61405
sha256=19e3992237906d1add92de6e691831cbb09c52f4e4b185f32c163555a980b07
d
  Stored in directory:
/root/.cache/pip/wheels/01/c0/af/77c1cf53a1be9e42a52b48e5af2169d40ec2e
89f7362489dd0
  Building wheel for antlr4-python3-runtime (setup.py) ... e:
filename=antlr4_python3_runtime-4.9.3-py3-none-any.whl size=144554
sha256=0a49064e8f234bb6260243dd510726907c4cee5959bb9594c03216486193d1c
1
  Stored in directory:
/root/.cache/pip/wheels/12/93/dd/1f6a127edc45659556564c5730f6d4e300888
f4bca2d4c5a88
Successfully built detectron2 fvcore antlr4-python3-runtime
Installing collected packages: antlr4-python3-runtime, yacs,
portalocker, pathspec, omegaconf, mpyc-extensions, iopath, hydra-core,
black, fvcore, detectron2
Successfully installed antlr4-python3-runtime-4.9.3 black-23.7.0
detectron2-0.6 fvcore-0.1.5.post20221221 hydra-core-1.3.2 iopath-0.1.9
```

```
mypy-extensions-1.0.0 omegaconf-2.3.0 pathspec-0.11.1 portalocker-2.7.0 yacs-0.1.8
```

```
!python -m pip install pyyaml==5.1
```

```
Collecting pyyaml==5.1
```

```
  Downloading PyYAML-5.1.tar.gz (274 kB)
```

```
0.0/274.2 kB ? eta -:--:--  
112.6/274.2 kB 3.4 MB/s eta  
0:00:01 274.2/274.2 kB 5.2
```

```
MB/s eta 0:00:00
```

```
etadate (setup.py) ... l
```

```
Building wheel for pyyaml (setup.py) ... l: filename=PyYAML-5.1-cp310-cp310-linux_x86_64.whl size=44090
```

```
sha256=4fbb5a68d2a7e7acbc6f818aaf6f79187097536b5ba78ac141fdebfe4786858
```

```
Stored in directory:
```

```
/root/.cache/pip/wheels/70/83/31/975b737609aba39a4099d471d5684141c1fdc3404f97e7f68a
```

```
Successfully built pyyaml
```

```
Installing collected packages: pyyaml
```

```
Attempting uninstall: pyyaml
```

```
Found existing installation: PyYAML 6.0.1
```

```
Uninstalling PyYAML-6.0.1:
```

```
Successfully uninstalled PyYAML-6.0.1
```

```
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
```

```
dask 2022.12.1 requires pyyaml>=5.3.1, but you have pyyaml 5.1 which is incompatible.
```

```
flax 0.7.0 requires PyYAML>=5.4.1, but you have pyyaml 5.1 which is incompatible.
```

```
Successfully installed pyyaml-5.1
```

```
import torch, detectron2
```

```
!nvcc --version
```

```
TORCH_VERSION = ".".join(torch.__version__.split(".")[0:2])
```

```
CUDA_VERSION = torch.__version__.split("+")[-1]
```

```
print("torch: ", TORCH_VERSION, "; cuda: ", CUDA_VERSION)
```

```
print("detectron2:", detectron2.__version__)
```

```
nvcc: NVIDIA (R) Cuda compiler driver
```

```
Copyright (c) 2005-2022 NVIDIA Corporation
```

```
Built on Wed_Sep_21_10:33:58_PDT_2022
```

```
Cuda compilation tools, release 11.8, V11.8.89
```

```
Build cuda_11.8.r11.8/compiler.31833905_0
```

```
torch: 2.0 ; cuda: cu118
```

```
detectron2: 0.6
```

```

import detectron2
from detectron2.utils.logger import setup_logger
setup_logger()

# import some common libraries
import numpy as np
import cv2
import matplotlib.pyplot as plt

# import some common detectron2 utilities
from detectron2 import model_zoo
from detectron2.engine import DefaultPredictor
from detectron2.config import get_cfg
from detectron2.utils.visualizer import Visualizer
from detectron2.data import MetadataCatalog, DatasetCatalog

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

import os
import numpy as np
import json
from detectron2.structures import BoxMode

def get_r_dicts(directory):
    classes = ['unmelted particle', 'porosity', 'microcrack']
    dataset_dicts = []
    for idx, filename in enumerate([file for file in
os.listdir(directory) if file.endswith('.json')]):
        json_file = os.path.join(directory, filename)
        with open(json_file) as f:
            img_anns = json.load(f)

        record = {}

        filename = os.path.join(directory, img_anns["imagePath"])

        record["file_name"] = filename
        record["image_id"] = idx
        record["height"] = 528
        record["width"] = 960

        annos = img_anns["shapes"]
        objs = []
        for anno in annos:
            px = [a[0] for a in anno['points']]
            py = [a[1] for a in anno['points']]
            poly = [(x, y) for x, y in zip(px, py)]

```

```



        poly = [p for x in poly for p in x]
        obj = {
            "bbox": [np.min(px), np.min(py), np.max(px),
np.max(py)],
            "bbox_mode": BoxMode.XYXY_ABS,
            "segmentation": [poly],
            "category_id": classes.index(anno['label']),
            "iscrowd": 0
        }
        objs.append(obj)
        record["annotations"] = objs
        dataset_dicts.append(record)
    return dataset_dicts


from detectron2.data import DatasetCatalog, MetadataCatalog
for d in ["train", "test"]:
    DatasetCatalog.register("p_" + d, lambda d=d:
get_r_dicts('/content/drive/MyDrive/Mahabub/' + d))
    MetadataCatalog.get("p_" + d).set(thing_classes=['unmelted
particle', 'porosity', 'microcrack'])
r_metadata = MetadataCatalog.get("p_train")

import random

dataset_dicts = get_r_dicts("/content/drive/MyDrive/Mahabub/train")
for d in random.sample(dataset_dicts, 3):
    img = cv2.imread(d["file_name"])
    v = Visualizer(img[:, :, ::-1], metadata=r_metadata, scale=0.5)
    v = v.draw_dataset_dict(d)
    plt.figure(figsize = (14, 10))
    plt.imshow(cv2.cvtColor(v.get_image()[:, :, ::-1],
cv2.COLOR_BGR2RGB))
    plt.show()

```





```
DatasetCatalog.remove("p_train")
DatasetCatalog.remove("p_test")
```

```
-----
-----
KeyError                                Traceback (most recent call
last)
```

```
<ipython-input-16-9f969f888a51> in <cell line: 1>()
```

```
----> 1 DatasetCatalog.remove("p_train")
```

```
      2 DatasetCatalog.remove("p_test")
```

```
/usr/local/lib/python3.10/dist-packages/detectron2/data/catalog.py in
remove(self, name)
```

```
    71     Alias of ``pop``.
```

```
    72     """
```

```
---> 73     self.pop(name)
```

```
    74
```

```
    75     def __str__(self):
```

```
/usr/lib/python3.10/_collections_abc.py in pop(self, key, default)
```

```
    955     """
```

```
    956     try:
```

```
--> 957         value = self[key]
```

```
    958     except KeyError:
```

```
    959         if default is self.__marker:
```

```
/usr/lib/python3.10/collections/_init_.py in __getitem__(self, key)
```

```
   1104     if hasattr(self.__class__, "__missing__"):
```

```

1105         return self.__class__.__missing__(self, key)
-> 1106         raise KeyError(key)
1107
1108     def __setitem__(self, key, item):

```

KeyError: 'p_train'

```

from detectron2.engine import DefaultTrainer
from detectron2.config import get_cfg
from detectron2.model_zoo import model_zoo

cfg = get_cfg()
cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/retinanet_R_101_FPN_3x.yaml"))
cfg.DATASETS.TRAIN = ("p_train",)
cfg.DATASETS.TEST = ()
cfg.DATALOADER.NUM_WORKERS = 2
cfg.MODEL.WEIGHTS =
model_zoo.get_checkpoint_url("COCO-Detection/retinanet_R_101_FPN_3x.yaml")
cfg.SOLVER.IMS_PER_BATCH = 2
cfg.SOLVER.BASE_LR = 0.00025
cfg.SOLVER.MAX_ITER = 2000
cfg.SOLVER.STEPS = [] # do not decay learning rate
cfg.MODEL.RETINANET.NUM_CLASSES = 3

os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
trainer = DefaultTrainer(cfg)
trainer.resume_or_load(resume=False)
trainer.train()

```

WARNING:fvcore.common.config:Loading config
/usr/local/lib/python3.10/dist-packages/detectron2/model_zoo/configs/
COCO-Detection/./Base-RetinaNet.yaml with yaml.unsafe_load. Your
machine may be at risk if the file contains malicious content.

[07/25 11:20:58 d2.engine.defaults]: Model:

```

RetinaNet(
  (backbone): FPN(
    (fpn_lateral3): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (fpn_lateral4): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (fpn_lateral5): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output5): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),

```

```

padding=(1, 1))
    (top_block): LastLevelP6P7(
        (p6): Conv2d(2048, 256, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1))
        (p7): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1))
    )
    (bottom_up): ResNet(
        (stem): BasicStem(
            (conv1): Conv2d(
                3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False
            )
            (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
        )
        (res2): Sequential(
            (0): BottleneckBlock(
                (shortcut): Conv2d(
                    64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
                )
                (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
            )
            (conv1): Conv2d(
                64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
            )
            (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
        )
            (conv2): Conv2d(
                64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
            )
            (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
        )
            (conv3): Conv2d(
                64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
            )
            (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
        )
            (1): BottleneckBlock(
                (conv1): Conv2d(
                    256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
                )
                (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
            )
            (conv2): Conv2d(
                64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
            )
            (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
        )
            (conv3): Conv2d(
                64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
            )
            (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
        )
    )

```

```

    )
    (2): BottleneckBlock(
      (conv1): Conv2d(
        256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      )
      (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False
        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      )
      (conv3): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
    )
  )
  (res3): Sequential(
    (0): BottleneckBlock(
      (shortcut): Conv2d(
        256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
      )
      (conv1): Conv2d(
        256, 128, kernel_size=(1, 1), stride=(2, 2), bias=False
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
      )
      (conv2): Conv2d(
        128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
      )
      (conv3): Conv2d(
        128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
      )
    )
    (1): BottleneckBlock(
      (conv1): Conv2d(
        512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
      )
      (conv2): Conv2d(
        128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
      )
      (conv3): Conv2d(
        128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False

```

```

        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
(2): BottleneckBlock(
  (conv1): Conv2d(
    512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv2): Conv2d(
    128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv3): Conv2d(
    128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
  )
)
(3): BottleneckBlock(
  (conv1): Conv2d(
    512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv2): Conv2d(
    128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
  (conv3): Conv2d(
    128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
  )
)
)
(res4): Sequential(
  (0): BottleneckBlock(
    (shortcut): Conv2d(
      512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
    (conv1): Conv2d(
      512, 256, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
  )
)

```

```

        (conv3): Conv2d(
          256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
          (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
        )
      )
    (1): BottleneckBlock(
      (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
      )
    )
    (2): BottleneckBlock(
      (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
      )
    )
    (3): BottleneckBlock(
      (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
      )
    )

```

```

    )
    (4): BottleneckBlock(
      (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
      )
    )
    (5): BottleneckBlock(
      (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
      )
    )
    (6): BottleneckBlock(
      (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
      )
    )
    (7): BottleneckBlock(
      (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False

```

```

        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
(8): BottleneckBlock(
    (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
(9): BottleneckBlock(
    (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
(10): BottleneckBlock(
    (conv1): Conv2d(
        1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,

```



```

1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (11): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (12): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (13): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(

```

```

        256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
)
(14): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(15): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(16): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
)

```

```

(17): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(18): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(19): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv2): Conv2d(
    256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
  )
  (conv3): Conv2d(
    256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
  )
)
(20): BottleneckBlock(
  (conv1): Conv2d(
    1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
    (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)

```

```

    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (21): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
  (22): BottleneckBlock(
    (conv1): Conv2d(
      1024, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv2): Conv2d(
      256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (conv3): Conv2d(
      256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False
      (norm): FrozenBatchNorm2d(num_features=1024, eps=1e-05)
    )
  )
)
(res5): Sequential(
  (0): BottleneckBlock(
    (shortcut): Conv2d(
      1024, 2048, kernel_size=(1, 1), stride=(2, 2), bias=False
      (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
    (conv1): Conv2d(
      1024, 512, kernel_size=(1, 1), stride=(2, 2), bias=False

```

```

        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv2): Conv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv3): Conv2d(
        512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
)
(1): BottleneckBlock(
    (conv1): Conv2d(
        2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv2): Conv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv3): Conv2d(
        512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
)
(2): BottleneckBlock(
    (conv1): Conv2d(
        2048, 512, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv2): Conv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False
        (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv3): Conv2d(
        512, 2048, kernel_size=(1, 1), stride=(1, 1), bias=False
        (norm): FrozenBatchNorm2d(num_features=2048, eps=1e-05)
    )
)
)
)
)
(head): RetinaNetHead(
    (cls_subnet): Sequential(
        (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
)

```

```

        (1): ReLU()
        (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (3): ReLU()
        (4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (5): ReLU()
        (6): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (7): ReLU()
    )
    (bbox_subnet): Sequential(
      (0): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (1): ReLU()
      (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (3): ReLU()
      (4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (5): ReLU()
      (6): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
      (7): ReLU()
    )
    (cls_score): Conv2d(256, 27, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (bbox_pred): Conv2d(256, 36, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (anchor_generator): DefaultAnchorGenerator(
    (cell_anchors): BufferList()
  )
)

```

[07/25 11:20:58 d2.data.build]: Removed 0 images with no usable annotations. 42 images left.

[07/25 11:20:58 d2.data.build]: Distribution of instances among all 3 categories:

category	#instances	category	#instances	category
unmelted pa..	639	porosity	67	
microcrack	9			
total	715			

[07/25 11:20:58 d2.data.dataset_mapper]: [DatasetMapper] Augmentations

```

used in training: [ResizeShortestEdge(short_edge_length=(640, 672,
704, 736, 768, 800), max_size=1333, sample_style='choice'),
RandomFlip()]
[07/25 11:20:58 d2.data.build]: Using training sampler TrainingSampler
[07/25 11:20:58 d2.data.common]: Serializing the dataset using: <class
'detectron2.data.common._TorchSerializedList'>
[07/25 11:20:58 d2.data.common]: Serializing 42 elements to byte
tensors and concatenating them all ...
[07/25 11:20:58 d2.data.common]: Serialized dataset takes 0.16 MiB
[07/25 11:20:58 d2.checkpoint.detection_checkpoint]:
[DetectionCheckpointer] Loading from
https://dl.fbaipublicfiles.com/detectron2/COCO-Detection/retinanet_R_1
01_FPN_3x/190397697/model_final_971ab9.pkl ...
model_final_971ab9.pkl: 228MB [00:04, 46.1MB/s]

WARNING:fvcore.common.checkpoint:Skip loading parameter
'head.cls_score.weight' to the model due to incompatible shapes: (720,
256, 3, 3) in the checkpoint but (27, 256, 3, 3) in the model! You
might want to double check if this is expected.
WARNING:fvcore.common.checkpoint:Skip loading parameter
'head.cls_score.bias' to the model due to incompatible shapes: (720,)
in the checkpoint but (27,) in the model! You might want to double
check if this is expected.
WARNING:fvcore.common.checkpoint:Some model parameters or buffers are
not found in the checkpoint:
head.cls_score.{bias, weight}
WARNING:fvcore.common.checkpoint:The checkpoint state_dict contains
keys that are not used by the model:
  pixel_mean
  pixel_std

[07/25 11:21:03 d2.engine.train_loop]: Starting training from
iteration 0

/usr/local/lib/python3.10/dist-packages/torch/functional.py:504:
UserWarning: torch.meshgrid: in an upcoming release, it will be
required to pass the indexing argument. (Triggered internally at
../aten/src/ATen/native/TensorShape.cpp:3483.)
  return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]

[07/25 11:21:17 d2.utils.events]: eta: 0:10:53 iter: 19 total_loss:
1.898 loss_cls: 1.337 loss_box_reg: 0.5127 time: 0.5389
last_time: 0.2966 data_time: 0.0591 last_data_time: 0.0160 lr:
4.9953e-06 max_mem: 2833M
[07/25 11:21:26 d2.utils.events]: eta: 0:09:42 iter: 39 total_loss:
1.342 loss_cls: 0.9546 loss_box_reg: 0.3698 time: 0.4002
last_time: 0.2556 data_time: 0.0075 last_data_time: 0.0091 lr:
9.9902e-06 max_mem: 2916M
[07/25 11:21:32 d2.utils.events]: eta: 0:09:39 iter: 59 total_loss:

```

```
1.214 loss_cls: 0.8363 loss_box_reg: 0.3802 time: 0.3696
last_time: 0.2571 data_time: 0.0172 last_data_time: 0.0183 lr:
1.4985e-05 max_mem: 2916M
[07/25 11:21:37 d2.utils.events]: eta: 0:09:20 iter: 79 total_loss:
1.089 loss_cls: 0.7315 loss_box_reg: 0.3571 time: 0.3419
last_time: 0.2470 data_time: 0.0082 last_data_time: 0.0061 lr:
1.998e-05 max_mem: 2916M
[07/25 11:21:43 d2.utils.events]: eta: 0:09:06 iter: 99 total_loss:
1.077 loss_cls: 0.6711 loss_box_reg: 0.3927 time: 0.3278
last_time: 0.2880 data_time: 0.0096 last_data_time: 0.0103 lr:
2.4975e-05 max_mem: 2916M
[07/25 11:21:49 d2.utils.events]: eta: 0:09:03 iter: 119
total_loss: 0.9542 loss_cls: 0.5712 loss_box_reg: 0.3491 time:
0.3238 last_time: 0.2899 data_time: 0.0138 last_data_time: 0.0080
lr: 2.997e-05 max_mem: 2916M
[07/25 11:21:54 d2.utils.events]: eta: 0:08:46 iter: 139
total_loss: 0.9094 loss_cls: 0.5334 loss_box_reg: 0.3735 time:
0.3144 last_time: 0.2410 data_time: 0.0093 last_data_time: 0.0066
lr: 3.4965e-05 max_mem: 2916M
[07/25 11:22:01 d2.utils.events]: eta: 0:08:47 iter: 159
total_loss: 0.8444 loss_cls: 0.4552 loss_box_reg: 0.3222 time:
0.3157 last_time: 0.3932 data_time: 0.0163 last_data_time: 0.0227
lr: 3.996e-05 max_mem: 2916M
[07/25 11:22:06 d2.utils.events]: eta: 0:08:36 iter: 179
total_loss: 0.7493 loss_cls: 0.4229 loss_box_reg: 0.3589 time:
0.3107 last_time: 0.2479 data_time: 0.0108 last_data_time: 0.0106
lr: 4.4955e-05 max_mem: 2916M
[07/25 11:22:11 d2.utils.events]: eta: 0:08:16 iter: 199
total_loss: 0.8247 loss_cls: 0.4436 loss_box_reg: 0.3443 time:
0.3057 last_time: 0.3068 data_time: 0.0079 last_data_time: 0.0255
lr: 4.995e-05 max_mem: 2916M
[07/25 11:22:17 d2.utils.events]: eta: 0:08:19 iter: 219
total_loss: 0.8362 loss_cls: 0.4395 loss_box_reg: 0.3844 time:
0.3051 last_time: 0.2436 data_time: 0.0135 last_data_time: 0.0075
lr: 5.4945e-05 max_mem: 2916M
[07/25 11:22:22 d2.utils.events]: eta: 0:07:59 iter: 239
total_loss: 0.7496 loss_cls: 0.4138 loss_box_reg: 0.3384 time:
0.2997 last_time: 0.2066 data_time: 0.0088 last_data_time: 0.0056
lr: 5.994e-05 max_mem: 2916M
[07/25 11:22:28 d2.utils.events]: eta: 0:07:47 iter: 259
total_loss: 0.6571 loss_cls: 0.3488 loss_box_reg: 0.313 time:
0.2976 last_time: 0.2917 data_time: 0.0093 last_data_time: 0.0058
lr: 6.4935e-05 max_mem: 2916M
[07/25 11:22:33 d2.utils.events]: eta: 0:07:46 iter: 279
total_loss: 0.6601 loss_cls: 0.3461 loss_box_reg: 0.3373 time:
0.2967 last_time: 0.3037 data_time: 0.0081 last_data_time: 0.0080
lr: 6.993e-05 max_mem: 2916M
[07/25 11:22:38 d2.utils.events]: eta: 0:07:31 iter: 299
total_loss: 0.652 loss_cls: 0.3337 loss_box_reg: 0.3176 time:
```


0.2933 last_time: 0.2124 data_time: 0.0096 last_data_time: 0.0091
lr: 7.4925e-05 max_mem: 2916M
[07/25 11:22:44 d2.utils.events]: eta: 0:07:28 iter: 319
total_loss: 0.7376 loss_cls: 0.3626 loss_box_reg: 0.3792 time:
0.2942 last_time: 0.3536 data_time: 0.0096 last_data_time: 0.0124
lr: 7.992e-05 max_mem: 2916M
[07/25 11:22:50 d2.utils.events]: eta: 0:07:23 iter: 339
total_loss: 0.6173 loss_cls: 0.2935 loss_box_reg: 0.3294 time:
0.2929 last_time: 0.2999 data_time: 0.0079 last_data_time: 0.0073
lr: 8.4915e-05 max_mem: 2916M
[07/25 11:22:55 d2.utils.events]: eta: 0:07:17 iter: 359
total_loss: 0.5881 loss_cls: 0.2813 loss_box_reg: 0.3242 time:
0.2916 last_time: 0.2426 data_time: 0.0067 last_data_time: 0.0057
lr: 8.991e-05 max_mem: 2916M
[07/25 11:23:02 d2.utils.events]: eta: 0:07:15 iter: 379
total_loss: 0.6232 loss_cls: 0.2795 loss_box_reg: 0.3251 time:
0.2928 last_time: 0.2487 data_time: 0.0183 last_data_time: 0.0053
lr: 9.4905e-05 max_mem: 2916M
[07/25 11:23:07 d2.utils.events]: eta: 0:07:11 iter: 399
total_loss: 0.5777 loss_cls: 0.2686 loss_box_reg: 0.3108 time:
0.2922 last_time: 0.3035 data_time: 0.0099 last_data_time: 0.0082
lr: 9.99e-05 max_mem: 2916M
[07/25 11:23:13 d2.utils.events]: eta: 0:07:04 iter: 419
total_loss: 0.6232 loss_cls: 0.2818 loss_box_reg: 0.3244 time:
0.2920 last_time: 0.2892 data_time: 0.0113 last_data_time: 0.0226
lr: 0.0001049 max_mem: 2916M
[07/25 11:23:19 d2.utils.events]: eta: 0:06:58 iter: 439
total_loss: 0.5455 loss_cls: 0.2584 loss_box_reg: 0.2989 time:
0.2915 last_time: 0.2955 data_time: 0.0091 last_data_time: 0.0103
lr: 0.00010989 max_mem: 2916M
[07/25 11:23:24 d2.utils.events]: eta: 0:06:52 iter: 459
total_loss: 0.5962 loss_cls: 0.2516 loss_box_reg: 0.3166 time:
0.2904 last_time: 0.2484 data_time: 0.0079 last_data_time: 0.0075
lr: 0.00011489 max_mem: 2916M
[07/25 11:23:30 d2.utils.events]: eta: 0:06:53 iter: 479
total_loss: 0.5636 loss_cls: 0.2575 loss_box_reg: 0.3079 time:
0.2914 last_time: 0.2468 data_time: 0.0171 last_data_time: 0.0076
lr: 0.00011988 max_mem: 2916M
[07/25 11:23:35 d2.utils.events]: eta: 0:06:41 iter: 499
total_loss: 0.5418 loss_cls: 0.2414 loss_box_reg: 0.3037 time:
0.2898 last_time: 0.2116 data_time: 0.0083 last_data_time: 0.0087
lr: 0.00012488 max_mem: 2916M
[07/25 11:23:41 d2.utils.events]: eta: 0:06:35 iter: 519
total_loss: 0.5535 loss_cls: 0.2746 loss_box_reg: 0.2808 time:
0.2895 last_time: 0.3319 data_time: 0.0094 last_data_time: 0.0229
lr: 0.00012987 max_mem: 2932M
[07/25 11:23:47 d2.utils.events]: eta: 0:06:32 iter: 539
total_loss: 0.5542 loss_cls: 0.2335 loss_box_reg: 0.3208 time:
0.2896 last_time: 0.2667 data_time: 0.0090 last_data_time: 0.0059

```
lr: 0.00013487 max_mem: 2932M
[07/25 11:23:52 d2.utils.events]: eta: 0:06:25 iter: 559
total_loss: 0.5331 loss_cls: 0.2238 loss_box_reg: 0.3102 time:
0.2888 last_time: 0.3183 data_time: 0.0079 last_data_time: 0.0089
lr: 0.00013986 max_mem: 2932M
[07/25 11:23:58 d2.utils.events]: eta: 0:06:24 iter: 579
total_loss: 0.5427 loss_cls: 0.2479 loss_box_reg: 0.279 time:
0.2896 last_time: 0.3561 data_time: 0.0160 last_data_time: 0.0068
lr: 0.00014486 max_mem: 2932M
[07/25 11:24:04 d2.utils.events]: eta: 0:06:19 iter: 599
total_loss: 0.5354 loss_cls: 0.2206 loss_box_reg: 0.308 time:
0.2891 last_time: 0.3023 data_time: 0.0100 last_data_time: 0.0060
lr: 0.00014985 max_mem: 2932M
[07/25 11:24:09 d2.utils.events]: eta: 0:06:10 iter: 619
total_loss: 0.492 loss_cls: 0.2141 loss_box_reg: 0.2719 time:
0.2883 last_time: 0.3891 data_time: 0.0082 last_data_time: 0.0119
lr: 0.00015485 max_mem: 2932M
[07/25 11:24:16 d2.utils.events]: eta: 0:06:08 iter: 639
total_loss: 0.543 loss_cls: 0.2301 loss_box_reg: 0.3034 time:
0.2892 last_time: 0.2080 data_time: 0.0129 last_data_time: 0.0070
lr: 0.00015984 max_mem: 2932M
[07/25 11:24:21 d2.utils.events]: eta: 0:05:59 iter: 659
total_loss: 0.4984 loss_cls: 0.195 loss_box_reg: 0.2852 time:
0.2884 last_time: 0.2529 data_time: 0.0098 last_data_time: 0.0074
lr: 0.00016484 max_mem: 2932M
[07/25 11:24:27 d2.utils.events]: eta: 0:05:57 iter: 679
total_loss: 0.5158 loss_cls: 0.2149 loss_box_reg: 0.3012 time:
0.2895 last_time: 0.2539 data_time: 0.0146 last_data_time: 0.0126
lr: 0.00016983 max_mem: 2932M
[07/25 11:24:33 d2.utils.events]: eta: 0:05:52 iter: 699
total_loss: 0.4627 loss_cls: 0.1968 loss_box_reg: 0.2675 time:
0.2892 last_time: 0.2998 data_time: 0.0088 last_data_time: 0.0092
lr: 0.00017483 max_mem: 2932M
[07/25 11:24:38 d2.utils.events]: eta: 0:05:42 iter: 719
total_loss: 0.5467 loss_cls: 0.2132 loss_box_reg: 0.3089 time:
0.2884 last_time: 0.2506 data_time: 0.0102 last_data_time: 0.0128
lr: 0.00017982 max_mem: 2932M
[07/25 11:24:44 d2.utils.events]: eta: 0:05:41 iter: 739
total_loss: 0.4644 loss_cls: 0.1969 loss_box_reg: 0.2673 time:
0.2891 last_time: 0.3119 data_time: 0.0161 last_data_time: 0.0074
lr: 0.00018482 max_mem: 2932M
[07/25 11:24:50 d2.utils.events]: eta: 0:05:34 iter: 759
total_loss: 0.4581 loss_cls: 0.1881 loss_box_reg: 0.269 time:
0.2886 last_time: 0.2601 data_time: 0.0104 last_data_time: 0.0072
lr: 0.00018981 max_mem: 2932M
[07/25 11:24:56 d2.utils.events]: eta: 0:05:30 iter: 779
total_loss: 0.5036 loss_cls: 0.1826 loss_box_reg: 0.2979 time:
0.2892 last_time: 0.3019 data_time: 0.0110 last_data_time: 0.0193
lr: 0.00019481 max_mem: 2932M
```

```
[07/25 11:25:02 d2.utils.events]: eta: 0:05:23 iter: 799
total_loss: 0.4101 loss_cls: 0.164 loss_box_reg: 0.242 time:
0.2887 last_time: 0.2241 data_time: 0.0093 last_data_time: 0.0068
lr: 0.0001998 max_mem: 2932M
[07/25 11:25:07 d2.utils.events]: eta: 0:05:15 iter: 819
total_loss: 0.4922 loss_cls: 0.1933 loss_box_reg: 0.2786 time:
0.2878 last_time: 0.2495 data_time: 0.0089 last_data_time: 0.0091
lr: 0.0002048 max_mem: 2932M
[07/25 11:25:13 d2.utils.events]: eta: 0:05:13 iter: 839
total_loss: 0.4722 loss_cls: 0.169 loss_box_reg: 0.2765 time:
0.2885 last_time: 0.2975 data_time: 0.0135 last_data_time: 0.0058
lr: 0.00020979 max_mem: 2932M
[07/25 11:25:18 d2.utils.events]: eta: 0:05:07 iter: 859
total_loss: 0.4291 loss_cls: 0.1609 loss_box_reg: 0.2591 time:
0.2880 last_time: 0.2497 data_time: 0.0097 last_data_time: 0.0082
lr: 0.00021479 max_mem: 2932M
[07/25 11:25:24 d2.utils.events]: eta: 0:05:03 iter: 879
total_loss: 0.4322 loss_cls: 0.1665 loss_box_reg: 0.2682 time:
0.2882 last_time: 0.2257 data_time: 0.0124 last_data_time: 0.0066
lr: 0.00021978 max_mem: 2932M
[07/25 11:25:30 d2.utils.events]: eta: 0:04:58 iter: 899
total_loss: 0.4773 loss_cls: 0.1971 loss_box_reg: 0.2855 time:
0.2883 last_time: 0.3308 data_time: 0.0082 last_data_time: 0.0055
lr: 0.00022478 max_mem: 2932M
[07/25 11:25:36 d2.utils.events]: eta: 0:04:54 iter: 919
total_loss: 0.3905 loss_cls: 0.1719 loss_box_reg: 0.2201 time:
0.2884 last_time: 0.3159 data_time: 0.0102 last_data_time: 0.0113
lr: 0.00022977 max_mem: 2932M
[07/25 11:25:43 d2.utils.events]: eta: 0:04:49 iter: 939
total_loss: 0.4373 loss_cls: 0.1593 loss_box_reg: 0.2642 time:
0.2893 last_time: 0.2525 data_time: 0.0149 last_data_time: 0.0067
lr: 0.00023477 max_mem: 2932M
[07/25 11:25:48 d2.utils.events]: eta: 0:04:43 iter: 959
total_loss: 0.4123 loss_cls: 0.1585 loss_box_reg: 0.2487 time:
0.2889 last_time: 0.2918 data_time: 0.0092 last_data_time: 0.0058
lr: 0.00023976 max_mem: 2932M
[07/25 11:25:54 d2.utils.events]: eta: 0:04:38 iter: 979
total_loss: 0.4576 loss_cls: 0.1672 loss_box_reg: 0.2805 time:
0.2892 last_time: 0.3940 data_time: 0.0073 last_data_time: 0.0066
lr: 0.00024476 max_mem: 2932M
[07/25 11:26:00 d2.utils.events]: eta: 0:04:32 iter: 999
total_loss: 0.3946 loss_cls: 0.1454 loss_box_reg: 0.2335 time:
0.2890 last_time: 0.2456 data_time: 0.0089 last_data_time: 0.0073
lr: 0.00024975 max_mem: 2932M
[07/25 11:26:05 d2.utils.events]: eta: 0:04:24 iter: 1019
total_loss: 0.4281 loss_cls: 0.1604 loss_box_reg: 0.2676 time:
0.2885 last_time: 0.3365 data_time: 0.0071 last_data_time: 0.0079
lr: 0.00025 max_mem: 2932M
[07/25 11:26:11 d2.utils.events]: eta: 0:04:20 iter: 1039
```

```
total_loss: 0.4316 loss_cls: 0.1485 loss_box_reg: 0.2804 time:
0.2891 last_time: 0.3082 data_time: 0.0081 last_data_time: 0.0051
lr: 0.00025 max_mem: 2932M
[07/25 11:26:17 d2.utils.events]: eta: 0:04:11 iter: 1059
total_loss: 0.4312 loss_cls: 0.1641 loss_box_reg: 0.2569 time:
0.2888 last_time: 0.2277 data_time: 0.0077 last_data_time: 0.0075
lr: 0.00025 max_mem: 2932M
[07/25 11:26:23 d2.utils.events]: eta: 0:04:09 iter: 1079
total_loss: 0.3822 loss_cls: 0.155 loss_box_reg: 0.2262 time:
0.2889 last_time: 0.2423 data_time: 0.0092 last_data_time: 0.0119
lr: 0.00025 max_mem: 2933M
[07/25 11:26:29 d2.utils.events]: eta: 0:04:04 iter: 1099
total_loss: 0.4115 loss_cls: 0.1493 loss_box_reg: 0.246 time:
0.2893 last_time: 0.3578 data_time: 0.0076 last_data_time: 0.0059
lr: 0.00025 max_mem: 2933M
[07/25 11:26:34 d2.utils.events]: eta: 0:03:56 iter: 1119
total_loss: 0.371 loss_cls: 0.1371 loss_box_reg: 0.2381 time:
0.2889 last_time: 0.1974 data_time: 0.0079 last_data_time: 0.0091
lr: 0.00025 max_mem: 2933M
[07/25 11:26:41 d2.utils.events]: eta: 0:03:53 iter: 1139
total_loss: 0.4015 loss_cls: 0.1447 loss_box_reg: 0.2621 time:
0.2894 last_time: 0.2518 data_time: 0.0169 last_data_time: 0.0067
lr: 0.00025 max_mem: 2933M
[07/25 11:26:46 d2.utils.events]: eta: 0:03:46 iter: 1159
total_loss: 0.4044 loss_cls: 0.1542 loss_box_reg: 0.2589 time:
0.2890 last_time: 0.2489 data_time: 0.0087 last_data_time: 0.0067
lr: 0.00025 max_mem: 2933M
[07/25 11:26:52 d2.utils.events]: eta: 0:03:42 iter: 1179
total_loss: 0.3583 loss_cls: 0.1276 loss_box_reg: 0.2335 time:
0.2892 last_time: 0.2796 data_time: 0.0106 last_data_time: 0.0142
lr: 0.00025 max_mem: 2933M
[07/25 11:26:58 d2.utils.events]: eta: 0:03:38 iter: 1199
total_loss: 0.402 loss_cls: 0.1527 loss_box_reg: 0.2519 time:
0.2892 last_time: 0.2015 data_time: 0.0099 last_data_time: 0.0077
lr: 0.00025 max_mem: 2933M
[07/25 11:27:03 d2.utils.events]: eta: 0:03:31 iter: 1219
total_loss: 0.3374 loss_cls: 0.1303 loss_box_reg: 0.213 time:
0.2889 last_time: 0.2983 data_time: 0.0080 last_data_time: 0.0082
lr: 0.00025 max_mem: 2933M
[07/25 11:27:09 d2.utils.events]: eta: 0:03:28 iter: 1239
total_loss: 0.3911 loss_cls: 0.1303 loss_box_reg: 0.247 time:
0.2891 last_time: 0.2958 data_time: 0.0148 last_data_time: 0.0248
lr: 0.00025 max_mem: 2933M
[07/25 11:27:15 d2.utils.events]: eta: 0:03:23 iter: 1259
total_loss: 0.3475 loss_cls: 0.1227 loss_box_reg: 0.2129 time:
0.2888 last_time: 0.2170 data_time: 0.0095 last_data_time: 0.0068
lr: 0.00025 max_mem: 2933M
[07/25 11:27:21 d2.utils.events]: eta: 0:03:18 iter: 1279
total_loss: 0.3463 loss_cls: 0.1254 loss_box_reg: 0.2277 time:
```

```
0.2888 last_time: 0.2898 data_time: 0.0088 last_data_time: 0.0193
lr: 0.00025 max_mem: 2933M
[07/25 11:27:27 d2.utils.events]: eta: 0:03:14 iter: 1299
total_loss: 0.375 loss_cls: 0.1275 loss_box_reg: 0.243 time:
0.2891 last_time: 0.2427 data_time: 0.0120 last_data_time: 0.0065
lr: 0.00025 max_mem: 2933M
[07/25 11:27:32 d2.utils.events]: eta: 0:03:08 iter: 1319
total_loss: 0.3391 loss_cls: 0.1131 loss_box_reg: 0.2216 time:
0.2888 last_time: 0.3229 data_time: 0.0084 last_data_time: 0.0060
lr: 0.00025 max_mem: 2933M
[07/25 11:27:38 d2.utils.events]: eta: 0:03:02 iter: 1339
total_loss: 0.3546 loss_cls: 0.1228 loss_box_reg: 0.2201 time:
0.2889 last_time: 0.3129 data_time: 0.0119 last_data_time: 0.0258
lr: 0.00025 max_mem: 2933M
[07/25 11:27:44 d2.utils.events]: eta: 0:02:56 iter: 1359
total_loss: 0.328 loss_cls: 0.12 loss_box_reg: 0.2137 time:
0.2888 last_time: 0.3208 data_time: 0.0090 last_data_time: 0.0070
lr: 0.00025 max_mem: 2933M
[07/25 11:27:49 d2.utils.events]: eta: 0:02:51 iter: 1379
total_loss: 0.3467 loss_cls: 0.1137 loss_box_reg: 0.2321 time:
0.2887 last_time: 0.3107 data_time: 0.0107 last_data_time: 0.0180
lr: 0.00025 max_mem: 2933M
[07/25 11:27:56 d2.utils.events]: eta: 0:02:45 iter: 1399
total_loss: 0.3683 loss_cls: 0.1364 loss_box_reg: 0.2283 time:
0.2890 last_time: 0.2216 data_time: 0.0083 last_data_time: 0.0072
lr: 0.00025 max_mem: 2933M
[07/25 11:28:01 d2.utils.events]: eta: 0:02:39 iter: 1419
total_loss: 0.3881 loss_cls: 0.1202 loss_box_reg: 0.2345 time:
0.2887 last_time: 0.2540 data_time: 0.0081 last_data_time: 0.0068
lr: 0.00025 max_mem: 2933M
[07/25 11:28:07 d2.utils.events]: eta: 0:02:34 iter: 1439
total_loss: 0.3022 loss_cls: 0.1028 loss_box_reg: 0.21 time:
0.2891 last_time: 0.4277 data_time: 0.0148 last_data_time: 0.0332
lr: 0.00025 max_mem: 2933M
[07/25 11:28:13 d2.utils.events]: eta: 0:02:28 iter: 1459
total_loss: 0.3418 loss_cls: 0.1177 loss_box_reg: 0.2241 time:
0.2888 last_time: 0.2504 data_time: 0.0087 last_data_time: 0.0054
lr: 0.00025 max_mem: 2933M
[07/25 11:28:18 d2.utils.events]: eta: 0:02:21 iter: 1479
total_loss: 0.3189 loss_cls: 0.1102 loss_box_reg: 0.2212 time:
0.2886 last_time: 0.2637 data_time: 0.0079 last_data_time: 0.0062
lr: 0.00025 max_mem: 2933M
[07/25 11:28:24 d2.utils.events]: eta: 0:02:18 iter: 1499
total_loss: 0.3248 loss_cls: 0.1065 loss_box_reg: 0.2139 time:
0.2889 last_time: 0.3436 data_time: 0.0107 last_data_time: 0.0062
lr: 0.00025 max_mem: 2933M
[07/25 11:28:30 d2.utils.events]: eta: 0:02:12 iter: 1519
total_loss: 0.2947 loss_cls: 0.09337 loss_box_reg: 0.1959 time:
0.2887 last_time: 0.3386 data_time: 0.0093 last_data_time: 0.0059
```

```
lr: 0.00025 max_mem: 2933M
[07/25 11:28:36 d2.utils.events]: eta: 0:02:07 iter: 1539
total_loss: 0.3202 loss_cls: 0.1131 loss_box_reg: 0.2114 time:
0.2890 last_time: 0.3257 data_time: 0.0129 last_data_time: 0.0306
lr: 0.00025 max_mem: 2933M
[07/25 11:28:42 d2.utils.events]: eta: 0:02:01 iter: 1559
total_loss: 0.2794 loss_cls: 0.09443 loss_box_reg: 0.1783 time:
0.2890 last_time: 0.2508 data_time: 0.0091 last_data_time: 0.0077
lr: 0.00025 max_mem: 2933M
[07/25 11:28:48 d2.utils.events]: eta: 0:01:55 iter: 1579
total_loss: 0.2951 loss_cls: 0.1025 loss_box_reg: 0.2068 time:
0.2889 last_time: 0.3952 data_time: 0.0097 last_data_time: 0.0060
lr: 0.00025 max_mem: 2933M
[07/25 11:28:54 d2.utils.events]: eta: 0:01:50 iter: 1599
total_loss: 0.293 loss_cls: 0.1096 loss_box_reg: 0.1861 time:
0.2892 last_time: 0.2447 data_time: 0.0125 last_data_time: 0.0084
lr: 0.00025 max_mem: 2933M
[07/25 11:28:59 d2.utils.events]: eta: 0:01:45 iter: 1619
total_loss: 0.3113 loss_cls: 0.09173 loss_box_reg: 0.2019 time:
0.2890 last_time: 0.2973 data_time: 0.0084 last_data_time: 0.0055
lr: 0.00025 max_mem: 2933M
[07/25 11:29:06 d2.utils.events]: eta: 0:01:39 iter: 1639
total_loss: 0.2679 loss_cls: 0.08808 loss_box_reg: 0.1852 time:
0.2894 last_time: 0.4436 data_time: 0.0132 last_data_time: 0.0306
lr: 0.00025 max_mem: 2933M
[07/25 11:29:11 d2.utils.events]: eta: 0:01:34 iter: 1659
total_loss: 0.335 loss_cls: 0.1045 loss_box_reg: 0.2239 time:
0.2892 last_time: 0.3349 data_time: 0.0082 last_data_time: 0.0051
lr: 0.00025 max_mem: 2933M
[07/25 11:29:17 d2.utils.events]: eta: 0:01:28 iter: 1679
total_loss: 0.3034 loss_cls: 0.09743 loss_box_reg: 0.2036 time:
0.2890 last_time: 0.2622 data_time: 0.0093 last_data_time: 0.0203
lr: 0.00025 max_mem: 2933M
[07/25 11:29:23 d2.utils.events]: eta: 0:01:22 iter: 1699
total_loss: 0.3534 loss_cls: 0.1181 loss_box_reg: 0.2276 time:
0.2895 last_time: 0.2460 data_time: 0.0127 last_data_time: 0.0075
lr: 0.00025 max_mem: 2933M
[07/25 11:29:28 d2.utils.events]: eta: 0:01:17 iter: 1719
total_loss: 0.308 loss_cls: 0.104 loss_box_reg: 0.1978 time:
0.2890 last_time: 0.3030 data_time: 0.0085 last_data_time: 0.0070
lr: 0.00025 max_mem: 2933M
[07/25 11:29:34 d2.utils.events]: eta: 0:01:10 iter: 1739
total_loss: 0.2836 loss_cls: 0.0934 loss_box_reg: 0.1862 time:
0.2891 last_time: 0.3960 data_time: 0.0116 last_data_time: 0.0059
lr: 0.00025 max_mem: 2933M
[07/25 11:29:41 d2.utils.events]: eta: 0:01:06 iter: 1759
total_loss: 0.2937 loss_cls: 0.08975 loss_box_reg: 0.2042 time:
0.2894 last_time: 0.3137 data_time: 0.0111 last_data_time: 0.0057
lr: 0.00025 max_mem: 2933M
```

```
[07/25 11:29:46 d2.utils.events]: eta: 0:01:00 iter: 1779
total_loss: 0.2816 loss_cls: 0.07974 loss_box_reg: 0.1917 time:
0.2893 last_time: 0.2762 data_time: 0.0075 last_data_time: 0.0089
lr: 0.00025 max_mem: 2933M
[07/25 11:29:53 d2.utils.events]: eta: 0:00:55 iter: 1799
total_loss: 0.2546 loss_cls: 0.08115 loss_box_reg: 0.1726 time:
0.2896 last_time: 0.2601 data_time: 0.0143 last_data_time: 0.0216
lr: 0.00025 max_mem: 2933M
[07/25 11:29:59 d2.utils.events]: eta: 0:00:50 iter: 1819
total_loss: 0.3138 loss_cls: 0.1011 loss_box_reg: 0.1969 time:
0.2897 last_time: 0.2998 data_time: 0.0094 last_data_time: 0.0069
lr: 0.00025 max_mem: 2933M
[07/25 11:30:05 d2.utils.events]: eta: 0:00:45 iter: 1839
total_loss: 0.2967 loss_cls: 0.09357 loss_box_reg: 0.2011 time:
0.2902 last_time: 0.2800 data_time: 0.0109 last_data_time: 0.0080
lr: 0.00025 max_mem: 2933M
[07/25 11:30:10 d2.utils.events]: eta: 0:00:39 iter: 1859
total_loss: 0.2646 loss_cls: 0.07796 loss_box_reg: 0.1805 time:
0.2899 last_time: 0.2939 data_time: 0.0072 last_data_time: 0.0052
lr: 0.00025 max_mem: 2933M
[07/25 11:30:16 d2.utils.events]: eta: 0:00:33 iter: 1879
total_loss: 0.2573 loss_cls: 0.07447 loss_box_reg: 0.1799 time:
0.2896 last_time: 0.4247 data_time: 0.0104 last_data_time: 0.0059
lr: 0.00025 max_mem: 2933M
[07/25 11:30:22 d2.utils.events]: eta: 0:00:27 iter: 1899
total_loss: 0.2752 loss_cls: 0.08372 loss_box_reg: 0.1788 time:
0.2898 last_time: 0.3119 data_time: 0.0167 last_data_time: 0.0083
lr: 0.00025 max_mem: 2933M
[07/25 11:30:28 d2.utils.events]: eta: 0:00:22 iter: 1919
total_loss: 0.2944 loss_cls: 0.09137 loss_box_reg: 0.2042 time:
0.2898 last_time: 0.2218 data_time: 0.0081 last_data_time: 0.0121
lr: 0.00025 max_mem: 2933M
[07/25 11:30:35 d2.utils.events]: eta: 0:00:16 iter: 1939
total_loss: 0.275 loss_cls: 0.08577 loss_box_reg: 0.1877 time:
0.2904 last_time: 0.2741 data_time: 0.0155 last_data_time: 0.0073
lr: 0.00025 max_mem: 2933M
[07/25 11:30:40 d2.utils.events]: eta: 0:00:11 iter: 1959
total_loss: 0.2537 loss_cls: 0.08225 loss_box_reg: 0.1816 time:
0.2903 last_time: 0.2557 data_time: 0.0084 last_data_time: 0.0096
lr: 0.00025 max_mem: 2933M
[07/25 11:30:46 d2.utils.events]: eta: 0:00:05 iter: 1979
total_loss: 0.2563 loss_cls: 0.08571 loss_box_reg: 0.1791 time:
0.2903 last_time: 0.3197 data_time: 0.0108 last_data_time: 0.0206
lr: 0.00025 max_mem: 2933M
[07/25 11:30:53 d2.utils.events]: eta: 0:00:00 iter: 1999
total_loss: 0.2599 loss_cls: 0.07856 loss_box_reg: 0.1832 time:
0.2904 last_time: 0.2115 data_time: 0.0096 last_data_time: 0.0068
lr: 0.00025 max_mem: 2933M
[07/25 11:30:54 d2.engine.hooks]: Overall training speed: 1998
iterations in 0:09:40 (0.2904 s / it)
```

```
[07/25 11:30:54 d2.engine.hooks]: Total training time: 0:09:46  
(0:00:05 on hooks)
```

```
# Look at training curves in tensorboard:
```


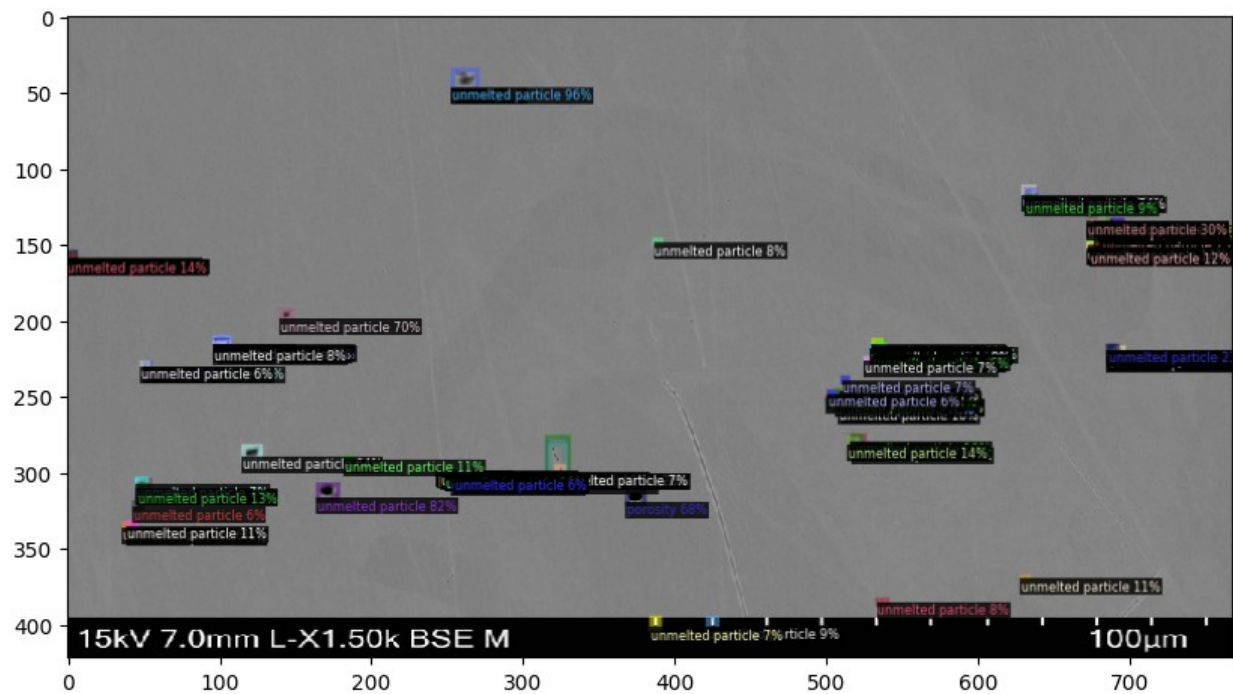
```
%load_ext tensorboard  
%tensorboard --logdir output
```

```
<IPython.core.display.Javascript object>
```

```
cfg.MODEL.WEIGHTS = os.path.join(cfg.OUTPUT_DIR, "model_final.pth")  
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.5  
cfg.DATASETS.TEST = ("p_test", )  
predictor = DefaultPredictor(cfg)
```

```
[07/25 11:31:48 d2.checkpoint.detection_checkpoint]:  
[DetectionCheckpointer] Loading from ./output/model_final.pth ...
```

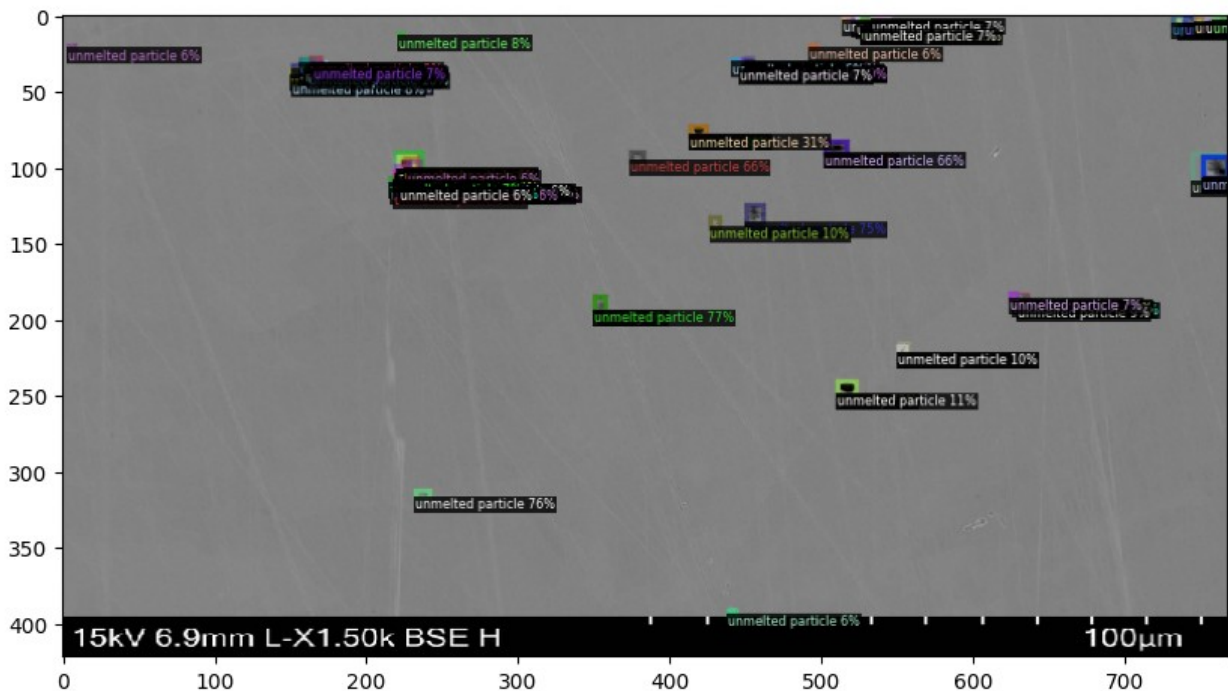
```
from detectron2.utils.visualizer import ColorMode  
dataset_dicts = get_r_dicts('/content/drive/MyDrive/Mahabub/train')  
for d in random.sample(dataset_dicts, 4):  
    im = cv2.imread(d["file_name"])  
    outputs = predictor(im)  
    v = Visualizer(im[:, :, ::-1],  
                  metadata=r_metadata,  
                  scale=0.8,  
                  instance_mode=ColorMode.IMAGE_BW # remove the  
colors of unsegmented pixels  
    )  
    v = v.draw_instance_predictions(outputs["instances"].to("cpu"))  
    plt.figure(figsize = (10, 10))  
    plt.imshow(cv2.cvtColor(v.get_image()[:, :, ::-1],  
cv2.COLOR_BGR2RGB))  
    plt.show()
```






```

        metadata=r_metadata,
        scale=0.8,
        instance_mode=ColorMode.IMAGE_BW # remove the
        colors of unsegmented pixels
    )
    v = v.draw_instance_predictions(outputs["instances"].to("cpu"))
    plt.figure(figsize = (10, 10))
    plt.imshow(cv2.cvtColor(v.get_image()[:, :, ::-1],
cv2.COLOR_BGR2RGB))
    plt.show()

```





```

from detectron2.evaluation import COCOEvaluator, inference_on_dataset
from detectron2.data import build_detection_test_loader
evaluator = COCOEvaluator("p_train", ['bbox'], False,
output_dir="./output/")
val_loader = build_detection_test_loader(cfg, "p_train")
print(inference_on_dataset(predictor.model, val_loader, evaluator))

[07/25 11:32:48 d2.evaluation.coco_evaluation]: Trying to convert
'p_train' to COCO format ...
[07/25 11:32:48 d2.data.datasets.coco]: Converting annotations of
dataset 'p_train' to COCO format ...
[07/25 11:32:48 d2.data.datasets.coco]: Converting dataset dicts into
COCO format
[07/25 11:32:49 d2.data.datasets.coco]: Conversion finished, #images:
42, #annotations: 715
[07/25 11:32:49 d2.data.datasets.coco]: Caching COCO format
annotations at './output/p_train_coco_format.json' ...
[07/25 11:32:49 d2.data.dataset_mapper]: [DatasetMapper] Augmentations
used in inference: [ResizeShortestEdge(short_edge_length=(800, 800),
max_size=1333, sample_style='choice')]
[07/25 11:32:49 d2.data.common]: Serializing the dataset using: <class
'detectron2.data.common._TorchSerializedList'>
[07/25 11:32:49 d2.data.common]: Serializing 42 elements to byte
tensors and concatenating them all ...
[07/25 11:32:49 d2.data.common]: Serialized dataset takes 0.16 MiB
[07/25 11:32:49 d2.evaluation.evaluator]: Start inference on 42
batches
[07/25 11:32:50 d2.evaluation.evaluator]: Inference done 11/42.

```



```
Dataloading: 0.0020 s/iter. Inference: 0.0564 s/iter. Eval: 0.0005
s/iter. Total: 0.0589 s/iter. ETA=0:00:01
[07/25 11:32:52 d2.evaluation.evaluator]: Total inference time:
0:00:02.265970 (0.061242 s / iter per device, on 1 devices)
[07/25 11:32:52 d2.evaluation.evaluator]: Total inference pure compute
time: 0:00:02 (0.055835 s / iter per device, on 1 devices)
[07/25 11:32:52 d2.evaluation.coco_evaluation]: Preparing results for
COCO format ...
[07/25 11:32:52 d2.evaluation.coco_evaluation]: Saving results to
./output/coco_instances_results.json
[07/25 11:32:52 d2.evaluation.coco_evaluation]: Evaluating predictions
with unofficial COCO API...
Loading and preparing results...
DONE (t=0.01s)
creating index...
index created!
[07/25 11:32:52 d2.evaluation.fast_eval_api]: Evaluate annotation type
*bbox*
[07/25 11:32:52 d2.evaluation.fast_eval_api]: COCOeval_opt.evaluate()
finished in 0.04 seconds.
[07/25 11:32:52 d2.evaluation.fast_eval_api]: Accumulating evaluation
results...
[07/25 11:32:52 d2.evaluation.fast_eval_api]:
COCOeval_opt.accumulate() finished in 0.02 seconds.
Average Precision (AP) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.539
Average Precision (AP) @[ IoU=0.50 | area= all |
maxDets=100 ] = 0.817
Average Precision (AP) @[ IoU=0.75 | area= all |
maxDets=100 ] = 0.554
Average Precision (AP) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.511
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.923
Average Precision (AP) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
1 ] = 0.276
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=
10 ] = 0.557
Average Recall (AR) @[ IoU=0.50:0.95 | area= all |
maxDets=100 ] = 0.634
Average Recall (AR) @[ IoU=0.50:0.95 | area= small |
maxDets=100 ] = 0.606
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium |
maxDets=100 ] = 0.988
Average Recall (AR) @[ IoU=0.50:0.95 | area= large |
maxDets=100 ] = -1.000
[07/25 11:32:52 d2.evaluation.coco_evaluation]: Evaluation results for
```

```

bbox:
|   AP   |   AP50   |   AP75   |   APs   |   APm   |   APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 53.878 | 81.661 | 55.377 | 51.118 | 92.266 | nan |
[07/25 11:32:52 d2.evaluation.coco_evaluation]: Some metrics cannot be
computed and is shown as NaN.
[07/25 11:32:52 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category          | AP      | category    | AP      | category    | AP
|
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
---|
| unmelted particle | 45.444 | porosity    | 66.404 | microcrack |
49.787 |
OrderedDict([('bbox', {'AP': 53.878198180479174, 'AP50':
81.66097785790757, 'AP75': 55.37691012964614, 'APs':
51.11774899305674, 'APm': 92.26622662266226, 'APl': nan, 'AP-unmelted
particle': 45.44378159428122, 'AP-porosity': 66.40377703606681, 'AP-
microcrack': 49.7870359110895}})])

```